

Call Stata from Python

Zhao Xu

Principal Software Engineer
StataCorp LLC

Introduction

Stata 16 introduces tight integration with Python allowing users to embed and execute Python code from within Stata. In this talk, I will demonstrate new functionality we have been working on: calling Stata from within Python. Note that this functionality is not available yet and is still a work in progress.

Contents

- [How it works](#)
- [Benefits](#)
- [Configuration and initialization](#)
- [Examples](#)
 - [Example 1: Basic usage](#)
 - [Example 2: Load dataset from Python](#)
 - [Example 3: Work with multiple datasets \(frames\)](#)
 - [Example 4: Integrate with Mata](#)
 - [Example 5: Call Stata using API functions](#)
- [Conclusion](#)

How it works

The Python package **pystata** provides two ways to interact with Stata:

- IPython magic commands
 - **%%stata** and **%stata**
 - **%%mata** and **%mata**
 - **%pystata**
- A suite of API functions
 - The **stata** module
 - The **config** module

The magic commands can be used to access Stata and Mata interactively in an IPython kernel-based environment.

- Jupyter Notebook/console
- Jupyter Lab/console
- Other environments that support the IPython kernel, such as Spyder IDE, PyCharm IDE, etc

The API functions can be used to interact with Stata and Mata in a command-line Python environment.

- Windows Command Prompt
- Unix terminal
- Python built-in IDLE, etc

The API functions can also be used together with the magic commands in the IPython environment. Both of them can be used with **Stata's Function Interface (sfi) module** to access Stata and Mata.

Benefits

In Python, with this integration, you can now:

- Use Stata's broad suite of estimation and post-estimation commands
 - Model estimations for various disciplines
 - Statistical inferences and predictions
 - Marginal effects and interaction analysis
 - Model specification, diagnostic, and goodness-of-fit analysis
- Create hundreds of thousands of publication-quality and distinctly styled graphics
- Make your research and work reproducible all the time using Stata's integrated version control
- Write and execute Stata and Python code in one environment
- Interact with each other by passing data and results back and forth
- And more...

Configuration and initialization

To get started, we need to configure the **pystata** package within Python so that it can be found and imported by Python. Suppose we have Stata installed in **C:/Program Files/Stata/**, it then can be initialized in Python as follows:

In [1]:

```
import sys
sys.path.append("C:/Program Files/Stata/utilities")
from pystata import config
config.init()
```

```

      _/ _/_/ _/ _/ _/ (R)
     _/_/ _/ _/_/ _/ _/ Stata Embedded
    Statistics/Data analysis
      MP - Parallel Edition
                                Copyright 1985-2019 StataCorp LLC
                                StataCorp
                                4905 Lakeway Drive
                                College Station, Texas 77845 USA
                                800-STATA-PC      https://www.stata.com
                                979-696-4600      stata@stata.com
                                979-696-4601 (fax)
```

Stata license: 10-user 4-core network perpetual

Serial number: 1

Licensed to: Stata Developer
StataCorp LLC

Notes:

1. Unicode is supported; see help `unicode_advice`.
2. More than 2 billion observations are allowed; see help `obs_advice`.
3. Maximum number of variables is set to 5,000; see help `set_maxvar`.

Examples

Example 1: Basic usage

To illustrate the general usage of calling Stata from Python, we use the automobile data. The data has mileage rating and weight of 74 automobiles. The variables of interest in the data are **mpg**, **weight**, and **foreign**. The **foreign** variable assumes the value 1 for foreign and 0 for domestic automobiles. We wish to analysis the relationship among the mileage rating, weight, and whether the automobile is foreign or domestic.

In [2]:

```
%%stata
use https://www.stata-press.com/data/r16/auto, clear
describe
```

```
. use https://www.stata-press.com/data/r16/auto, clear
(1978 Automobile Data)
```

```
. describe
```

```
Contains data from https://www.stata-press.com/data/r16/auto.dta
   obs:                74                1978 Automobile Data
   vars:                12                13 Apr 2018 17:45
                                   (_dta has notes)
```

variable name	storage type	display format	value label	variable label
make	str18	%-18s		Make and Model
price	int	%8.0gc		Price
mpg	int	%8.0g		Mileage (mpg)
rep78	int	%8.0g		Repair Record 1978
headroom	float	%6.1f		Headroom (in.)
trunk	int	%8.0g		Trunk space (cu. ft.)
weight	int	%8.0gc		Weight (lbs.)
length	int	%8.0g		Length (in.)
turn	int	%8.0g		Turn Circle (ft.)
displacement	int	%8.0g		Displacement (cu. in.)
gear_ratio	float	%6.2f		Gear Ratio
foreign	byte	%8.0g	origin	Car type

```
Sorted by: foreign
```

Then we obtain summaries of **mpg** and **weight** for the foreign and domestic cars.

In [3]:

```
%stata by foreign: summarize mpg weight
```

```
-> foreign = Domestic
```

Variable	Obs	Mean	Std. Dev.	Min	Max
mpg	52	19.82692	4.743297	12	34
weight	52	3317.115	695.3637	1800	4840


```
-> foreign = Foreign
```

Variable	Obs	Mean	Std. Dev.	Min	Max
mpg	22	24.77273	6.611187	14	41
weight	22	2315.909	433.0035	1760	3420

We visualize **mpg** and **weight** for each group of cars using the scatter plot.

In [4]:

```
%%stata
scatter mpg weight, by(foreign, total)
```



Then we fit a linear regression model of **mpg** on **weight** and **foreign**.

In [5]:

```
%%stata
regress mpg weight i.foreign
```

Source	SS	df	MS	Number of obs	=	74
Model	1619.2877	2	809.643849	F(2, 71)	=	69.75
Residual	824.171761	71	11.608053	Prob > F	=	0.0000
				R-squared	=	0.6627
				Adj R-squared	=	0.6532
Total	2443.45946	73	33.4720474	Root MSE	=	3.4071

mpg	Coef.	Std. Err.	t	P> t	[95% Conf. Interval]
weight	-.0065879	.0006371	-10.34	0.000	-.0078583 -.0053175
foreign					
Foreign	-1.650029	1.075994	-1.53	0.130	-3.7955 .4954422
_cons	41.6797	2.165547	19.25	0.000	37.36172 45.99768

Next, we use **margins** to calculate the mean predicted values for various values of **weight** in increments of 1,000 between 2,000 and 5,000 and each group of cars. We then use **marginsplot** to show the results graphically.

In [6]:

```
%%stata  
margins, at(weight=(2000(1000)5000)) over(foreign)  
marginsplot, by(foreign) xlabel(, angle(forty_five))
```

```
. margins, at(weight=(2000(1000)5000)) over(foreign)
```

```
Predictive margins                                Number of obs    =           74
```

```
Model VCE      : OLS
```

```
Expression    : Linear prediction, predict()
```

```
over          : foreign
```

```
1._at        : 0. foreign
               weight      =      2000
               1. foreign
               weight      =      2000

2._at        : 0. foreign
               weight      =      3000
               1. foreign
               weight      =      3000


3._at        : 0. foreign
               weight      =      4000
               1. foreign
               weight      =      4000

4._at        : 0. foreign
               weight      =      5000
               1. foreign
               weight      =      5000
```

	Delta-method				[95% Conf. Interval]	
	Margin	Std. Err.	t	P> t		
_at#foreign						
1#Domestic	28.50393	.9630195	29.60	0.000	26.58372	30.42414
1#Foreign	26.8539	.7537561	35.63	0.000	25.35095	28.35685
2#Domestic	21.91604	.5138592	42.65	0.000	20.89144	22.94065
2#Foreign	20.26601	.8471116	23.92	0.000	18.57692	21.95511
3#Domestic	15.32816	.6422785	23.87	0.000	14.04749	16.60882
3#Foreign	13.67813	1.295714	10.56	0.000	11.09455	16.26171
4#Domestic	8.74027	1.171673	7.46	0.000	6.404021	11.07652
4#Foreign	7.090241	1.857949	3.82	0.000	3.385596	10.79489

```
. marginsplot, by(foreign) xlabel(, angle(forty_five))
```

Variables that uniquely identify margins: weight foreign



Last, we create a partial-regression leverage plot for all the regressors using the **avplots** command.

In [7]:



```

%%stata
set scheme plottigblind
avplots

```

```
. set scheme plottigblind
```

```
. avplots
```



Example 2: Load dataset from Python

There are many ways to load data from Python into Stata's current dataset in memory. For example

1. Pandas dataframes and Numpy arrays can be loaded directly into Stata.
2. The **Data** and **Frame** classes within the **Stata Function Interface (sfi)** module provide multiple methods for loading data from Python.
3. Stata can read in data from a variety of sources, many of which can be created in Python: Excel files, CSV files, SPSS and SAS dataset, and various databases.

We have data from the **National Longitudinal Survey** on young women's wages reported from 1968 through 1988. This dataset is stored in a csv file named **nlswork.csv**.

The goal is to use Stata to fit a model of wage as a function of each woman's age, job tenure, birth year, and level of education. We believe that random shocks that affect a woman's wage also affect her job tenure, so we treat tenure as endogenous. As additional instruments, we use her union status, number of weeks worked in the past year, and a dummy indicating whether she lives in a metropolitan area.

The plan is to load the data using Pandas dataframe and fit a single-equation instrumental-variables regression via the **ivregress** command.

In [8]:

```
import pandas as pd
nlswork = pd.read_csv('nlswork.csv')
nlswork
```

Out[8]:

	idcode	year	birth_yr	age	race	msp	nev_mar	grade	collgrad	not_smsa	...	so
0	1	70	51	18.0	black	0.0	1.0	12.0	0	0.0	...	
1	1	71	51	19.0	black	1.0	0.0	12.0	0	0.0	...	
2	1	72	51	20.0	black	1.0	0.0	12.0	0	0.0	...	
3	1	73	51	21.0	black	1.0	0.0	12.0	0	0.0	...	
4	1	75	51	23.0	black	1.0	0.0	12.0	0	0.0	...	
...	
28529	5159	80	44	35.0	black	0.0	0.0	12.0	0	0.0	...	
28530	5159	82	44	37.0	black	0.0	0.0	12.0	0	0.0	...	
28531	5159	83	44	38.0	black	0.0	0.0	12.0	0	0.0	...	
28532	5159	85	44	40.0	black	0.0	0.0	12.0	0	0.0	...	
28533	5159	88	44	43.0	black	0.0	0.0	12.0	0	0.0	...	

28534 rows × 21 columns



Then we load the dataframe into Stata as current dataset and specify the labels to the variables of interest within Stata.

In [9]:

```
%%stata -d nlswork -force
label variable ln_wage "ln(wage/GNP deflator)"
label variable age "age in current year"
label variable birth_yr "birth year"
label variable grade "current grade completed"
label variable tenure "job tenure, in years"
label variable union "weeks unemployed last year"
label variable wks_work "weeks worked last year"
label variable msp "1 if married, spouse present"

describe
```

```
. label variable ln_wage "ln(wage/GNP deflator)"
. label variable age "age in current year"
. label variable birth_yr "birth year"
. label variable grade "current grade completed"
. label variable tenure "job tenure, in years"
. label variable union "weeks unemployed last year"
. label variable wks_work "weeks worked last year"
. label variable msp "1 if married, spouse present"
.
. describe
```

Contains data

```
obs:      28,534
vars:      21
```

variable name	storage type	display format	value label	variable label
idcode	long	%12.0g		
year	long	%12.0g		
birth_yr	long	%12.0g		birth year
age	double	%10.0g		age in current year
race	str9	%9s		
msp	double	%10.0g		1 if married, spouse present
nev_mar	double	%10.0g		
grade	double	%10.0g		current grade completed
collgrad	long	%12.0g		
not_smsa	double	%10.0g		
c_city	double	%10.0g		
south	double	%10.0g		
ind_code	double	%10.0g		
occ_code	double	%10.0g		
union	double	%10.0g		weeks unemployed last year
wks_ue	double	%10.0g		
ttl_exp	double	%10.0g		
tenure	double	%10.0g		job tenure, in years
hours	double	%10.0g		
wks_work	double	%10.0g		weeks worked last year
ln_wage	double	%10.0g		ln(wage/GNP deflator)

Sorted by:

Note: Dataset has changed since last saved.

Next, we fit the model and push Stata's estimation results into Python, such as the coefficient vector **e(b)** and variance-covariance matrix **e(V)**. The estimation results is stored in **steret**, which is a Python dictionary.

In [10]:

```
%%stata -eret steret
// fit the model using the gmm estimator
ivregress gmm ln_wage age c.age#c.age birth_yr grade ///
    (tenure = union wks_work msp), wmatrix(cluster idcode)

// e() results
ereturn list
```

```
. // fit the model using the gmm estimator
. ivregress gmm ln_wage age c.age#c.age birth_yr grade ///
> (tenure = union wks_work msp), wmatrix(cluster idcode)
```

```
Instrumental variables (GMM) regression      Number of obs   =    18,625
                                           Wald chi2(5)    =    1807.17
                                           Prob > chi2     =     0.0000
                                           R-squared       =          .
GMM weight matrix: Cluster (idcode)        Root MSE        =     .46951
```

(Std. Err. adjusted for 4,110 clusters in idcode)

ln_wage	Coef.	Robust Std. Err.	z	P> z	[95% Conf. Interval]	
tenure	.099221	.0037764	26.27	0.000	.0918194	.1066227
age	.0171146	.0066895	2.56	0.011	.0040034	.0302259
c.age#c.age	-.0005191	.000111	-4.68	0.000	-.0007366	-.0003016
birth_yr	-.0085994	.0021932	-3.92	0.000	-.012898	-.0043008
grade	.071574	.0029938	23.91	0.000	.0657062	.0774417
_cons	.8575071	.1616274	5.31	0.000	.5407231	1.174291

Instrumented: tenure

Instruments: age c.age#c.age birth_yr grade union wks_work msp

```
.
. // e() results
. ereturn list
```

scalars:

```
      e(J) = 11.88787679472382
      e(rss) = 4105.610892967217
      e(N) = 18625
      e(df_m) = 5
      e(rmse) = .4695055741800723
      e(mss) = -63.29808734918015
      e(r2) = .
      e(r2_a) = .
      e(chi2) = 1807.171481540709
      e(iterations) = 1
      e(N_clust) = 4110
      e(rank) = 6
```

macros:

```
      e(cmdline) : "ivregress gmm ln_wage age c.age#c.age birth_yr gr.."
      e(cmd) : "ivregress"
      e(estat_cmd) : "ivregress_estat"
      e(wmatrix) : "cluster idcode"
      e(vce) : "cluster"
      e(clustvar) : "idcode"
      e(vcetype) : "Robust"
      e(estimator) : "gmm"
      e(footnote) : "ivreg_footnote"
      e(marginsnotok) : "Residuals SCores"
      e(marginsok) : "XB default"
      e(predict) : "ivregress_p"
      e(depvar) : "ln_wage"
      e(exogr) : "age c.age#c.age birth_yr grade"
      e(insts) : "age c.age#c.age birth_yr grade union wks_work msp"
```

```
e(instd) : "tenure"  
e(title) : "Instrumental variables (GMM) regression"  
e(properties) : "b V"
```

matrices:

```
e(b) : 1 x 6  
e(V) : 6 x 6  
e(S) : 8 x 8  
e(W) : 8 x 8
```

functions:

```
e(sample)
```

Below, we list the entire contents of the Python dictionary **steret**.

In [11]:

```
steret
```


Out[11]:

```

{'e(J)': 11.887876794723823,
'e(rss)': 4105.610892967217,
'e(N)': 18625.0,
'e(df_m)': 5.0,
'e(rmse)': 0.46950557418007227,
'e(mss)': -63.29808734918015,
'e(r2)': 8.98846567431158e+307,
'e(r2_a)': 8.98846567431158e+307,
'e(chi2)': 1807.1714815407095,
'e(iterations)': 1.0,
'e(N_clust)': 4110.0,
'e(rank)': 6.0,
'e(cmdline)': 'ivregress gmm ln_wage age c.age#c.age birth_yr grade      (tenure
= union wks_work msp), wmatrix(cluster idcode)',
'e(cmd)': 'ivregress',
'e(estat_cmd)': 'ivregress_estat',
'e(robust_epilog)': 'ivregress_epilog',
'e(robust_prolog)': 'ivregress_prolog',
'e(wmatrix)': 'cluster idcode',
'e(vce)': 'cluster',
'e(clustvar)': 'idcode',
'e(vcetype)': 'Robust',
'e(estimator)': 'gmm',
'e(footnote)': 'ivreg_footnote',
'e(marginsprop)': 'nlinearize',
'e(marginsnotok)': 'Residuals SCores',
'e(marginsok)': 'XB default',
'e(predict)': 'ivregress_p',
'e(depvar)': 'ln_wage',
'e(exogr)': 'age c.age#c.age birth_yr grade',
'e(insts)': 'age c.age#c.age birth_yr grade union wks_work msp',
'e(instd)': 'tenure',
'e(title)': 'Instrumental variables (GMM) regression',
'e(properties)': 'b V',
'e(b)': array([[ 9.92210073e-02,  1.71146216e-02, -5.19104153e-04,
-8.59936559e-03,  7.15739528e-02,  8.57507064e-01]]),
'e(V)': array([[ 1.42613627e-05, -9.76241713e-07, -2.63864849e-08,
-7.77159613e-07, -3.01563638e-06,  8.56749367e-05],
[-9.76241713e-07,  4.47498135e-05, -7.33276730e-07,
 2.72940218e-06, -1.83647532e-06, -7.58296139e-04],
[-2.63864849e-08, -7.33276730e-07,  1.23108900e-08,
-3.57163864e-08,  3.86949853e-08,  1.17566367e-05],
[-7.77159613e-07,  2.72940218e-06, -3.57163864e-08,
 4.81015453e-06,  1.16191117e-07, -2.80185164e-04],
[-3.01563638e-06, -1.83647532e-06,  3.86949853e-08,
 1.16191117e-07,  8.96286674e-06, -9.00370663e-05],
[ 8.56749367e-05, -7.58296139e-04,  1.17566367e-05,
-2.80185164e-04, -9.00370663e-05,  2.61234292e-02]]),
'e(S)': array([[1.84225991e+07,  6.48996714e+08,  2.54724104e+07,  7.01502864e+06,
 5.36806324e+05,  1.53398668e+05,  3.94294503e+07,  3.22229084e+05],
[6.48996714e+08,  2.31708366e+10,  8.79392757e+08,  2.44144659e+08,
 1.86264122e+07,  5.37264373e+06,  1.38258192e+09,  1.12612714e+07],
[2.54724104e+07,  8.79392757e+08,  3.63515621e+07,  9.88420993e+06,
 7.59073876e+05,  2.14187357e+05,  5.50583611e+07,  4.50186699e+05],
[7.01502864e+06,  2.44144659e+08,  9.88420993e+06,  2.78761599e+06,
 2.07158962e+05,  5.83301079e+04,  1.51171391e+07,  1.23720912e+05],
[5.36806324e+05,  1.86264122e+07,  7.59073876e+05,  2.07158962e+05,
 1.59149056e+04,  4.49290588e+03,  1.15376185e+06,  9.47072185e+03],
[1.53398668e+05,  5.37264373e+06,  2.14187357e+05,  5.83301079e+04,

```

```

4. 49290588e+03, 3. 80211614e+03, 3. 26881494e+05, 2. 51143324e+03],
[3. 94294503e+07, 1. 38258192e+09, 5. 50583611e+07, 1. 51171391e+07,
1. 15376185e+06, 3. 26881494e+05, 8. 75368616e+07, 6. 78521312e+05],
[3. 22229084e+05, 1. 12612714e+07, 4. 50186699e+05, 1. 23720912e+05,
9. 47072185e+03, 2. 51143324e+03, 6. 78521312e+05, 8. 64778234e+03]]),
'e(W)': array([[ 7. 64077092e+00, -1. 16467817e-01, -3. 45417830e-01,
-8. 73750534e-02, -1. 00703558e+02, -3. 57303530e-01,
-3. 75313140e-02, -4. 81605062e-01],
[-1. 16467817e-01, 1. 81080290e-03, 7. 50889019e-03,
9. 15320423e-04, 1. 41514893e+00, 3. 91377503e-03,
2. 69838887e-04, 5. 60894455e-03],
[-3. 45417830e-01, 7. 50889019e-03, 2. 77106682e-01,
-2. 59628258e-02, -8. 66908704e+00, -6. 90985307e-02,
-1. 81124339e-02, -3. 10466629e-02],
[-8. 73750534e-02, 9. 15320423e-04, -2. 59628258e-02,
2. 03171641e-01, 5. 43785072e-01, 2. 37092939e-02,
-9. 75767432e-04, -1. 57125159e-02],
[-1. 00703558e+02, 1. 41514893e+00, -8. 66908704e+00,
5. 43785072e-01, 2. 05170227e+03, 7. 23678800e+00,
1. 26215920e+00, 5. 53726790e+00],
[-3. 57303530e-01, 3. 91377503e-03, -6. 90985307e-02,
2. 37092939e-02, 7. 23678800e+00, 6. 94839823e+00,
1. 33578947e-02, 4. 59819858e-01],
[-3. 75313140e-02, 2. 69838887e-04, -1. 81124339e-02,
-9. 75767432e-04, 1. 26215920e+00, 1. 33578947e-02,
7. 48017149e-03, 2. 96991881e-02],
[-4. 81605062e-01, 5. 60894455e-03, -3. 10466629e-02,
-1. 57125159e-02, 5. 53726790e+00, 4. 59819858e-01,
2. 96991881e-02, 5. 96402115e+00]]])}

```

You can also access specific elements of the dictionary. For example, you can access **e(b)** and **e(V)** by typing **steret['e(b)']** and **steret['e(V)']** in Python.

In [12]:

```
steret['e(V)']
```

Out[12]:

```

array([[ 1. 42613627e-05, -9. 76241713e-07, -2. 63864849e-08,
-7. 77159613e-07, -3. 01563638e-06, 8. 56749367e-05],
[-9. 76241713e-07, 4. 47498135e-05, -7. 33276730e-07,
2. 72940218e-06, -1. 83647532e-06, -7. 58296139e-04],
[-2. 63864849e-08, -7. 33276730e-07, 1. 23108900e-08,
-3. 57163864e-08, 3. 86949853e-08, 1. 17566367e-05],
[-7. 77159613e-07, 2. 72940218e-06, -3. 57163864e-08,
4. 81015453e-06, 1. 16191117e-07, -2. 80185164e-04],
[-3. 01563638e-06, -1. 83647532e-06, 3. 86949853e-08,
1. 16191117e-07, 8. 96286674e-06, -9. 00370663e-05],
[ 8. 56749367e-05, -7. 58296139e-04, 1. 17566367e-05,
-2. 80185164e-04, -9. 00370663e-05, 2. 61234292e-02]])

```

You can also access the above matrix using the **get()** method of the **Matrix** class in the **sfi** module.

In [13]:

```
from sfi import Matrix
import numpy as np

np.array(Matrix.get('e(V)'))
```

Out[13]:

```
array([[ 1.42613627e-05, -9.76241713e-07, -2.63864849e-08,
        -7.77159613e-07, -3.01563638e-06,  8.56749367e-05],
       [-9.76241713e-07,  4.47498135e-05, -7.33276730e-07,
        2.72940218e-06, -1.83647532e-06, -7.58296139e-04],
       [-2.63864849e-08, -7.33276730e-07,  1.23108900e-08,
        -3.57163864e-08,  3.86949853e-08,  1.17566367e-05],
       [-7.77159613e-07,  2.72940218e-06, -3.57163864e-08,
        4.81015453e-06,  1.16191117e-07, -2.80185164e-04],
       [-3.01563638e-06, -1.83647532e-06,  3.86949853e-08,
        1.16191117e-07,  8.96286674e-06, -9.00370663e-05],
       [ 8.56749367e-05, -7.58296139e-04,  1.17566367e-05,
        -2.80185164e-04, -9.00370663e-05,  2.61234292e-02]])
```

Example 3: Work with multiple datasets (frames)

Stata 16 introduced frames, allowing you to simultaneously work with multiple datasets in memory. The following example illustrates simple usage of multiple frames in Stata and how to switch frames between Stata and Python.

First, we load the *iris* dataset in Stata. This dataset is used in Fisher's (1936) article. Fisher obtained the iris data from Anderson (1935). The data consist of four features measured on 50 samples from each of three Iris species.

In [14]:

```
%%stata
use http://www.stata-press.com/data/r16/iris, clear
describe
label list species
```

```
. use http://www.stata-press.com/data/r16/iris, clear
(Iris data)
```

```
. describe
```

Contains data from http://www.stata-press.com/data/r16/iris.dta

```
obs:          150          Iris data
vars:          5          18 Jan 2018 13:23
                        (_dta has notes)
```

variable name	storage type	display format	value label	variable label
iris	byte	%10.0g	species	Iris species
seplen	double	%4.1f		Sepal length in cm
sepwid	double	%4.1f		Sepal width in cm
petlen	double	%4.1f		Petal length in cm
petwid	double	%4.1f		Petal width in cm

Sorted by:

```
. label list species
```

```
species:
      1 setosa
      2 versicolor
      3 virginica
```

Our goal is to build a classifier using those four features to detect the Iris type. We will use the **Random Forest** classification model within the **scikit-learn** Python package to achieve this goal. First, we split the data into the training and test datasets, and store the datasets in the **training** and **test** frames in Stata. The training dataset contains 80% of the observations and the test dataset contains 20% of the observations. Then we store the two frames into Python as Pandas dataframes with the same name.

In [15]:

```
%%stata -fouta training, test
// Split the original dataset into training and test
// dataset which contains 80% and 20% of observations respectively
splittsample, generate(svar, replace) split(0.8 0.2) show rseed(16)

// create two frames holding the the two datasets
frame put iris seplen sepwid petlen petwid if svar==1, into(training)
frame put iris seplen sepwid petlen petwid if svar==2, into(test)
```

```
. // Split the original dataset into training and test
. // dataset which contains 80% and 20% of observations respectively
. splittsample, generate(svar, replace) split(0.8 0.2) show rseed(16)
```

sva	Freq.	Percent	Cum.
1	120	80.00	80.00
2	30	20.00	100.00
Total	150	100.00	

```
.
. // create two frames holding the the two datasets
. frame put iris seplen sepwid petlen petwid if svar==1, into(training)

. frame put iris seplen sepwid petlen petwid if svar==2, into(test)
```

Now we have 2 NumPy arrays in Python, **training** and **test**. Below we split each array into two sub-arrays to store the features and labels separately.

In [16]:

```
X_train = training[:, 1:]
y_train = training[:, 0]
X_test = test[:, 1:]
y_test = test[:, 0]
```

Then we use **X_train** and **y_train** to train the classification model.

In [17]:

```
from sklearn.ensemble import RandomForestClassifier

clf = RandomForestClassifier(max_depth=2, random_state=16)
clf.fit(X_train, y_train)
```

Out[17]:

```
RandomForestClassifier(bootstrap=True, ccp_alpha=0.0, class_weight=None,
                       criterion='gini', max_depth=2, max_features='auto',
                       max_leaf_nodes=None, max_samples=None,
                       min_impurity_decrease=0.0, min_impurity_split=None,
                       min_samples_leaf=1, min_samples_split=2,
                       min_weight_fraction_leaf=0.0, n_estimators=100,
                       n_jobs=None, oob_score=False, random_state=16, verbose=0,
                       warm_start=False)
```

Next we use **X_test** and **y_test** to evaluate the performance of the training model. We also predict the species type of each flower and the probabilities that it belongs to the three species in the test dataset.

In [18]:

```
from sklearn import metrics

y_pred = clf.predict(X_test)
y_pred_prob = clf.predict_proba(X_test)
```

Next, in the **test** frame, we create a Byte variable **irispr** to store the predicted species types and three float variables to store the probabilities that each flower belongs to the three species types from the array **y_pred_prob**.

In [19]:

```
from sfi import Frame

fr = Frame.connect('test')
fr.addVarByte('irispr')
fr.addVarFloat('pr1')
fr.addVarFloat('pr2')
fr.addVarFloat('pr3')

fr.store('irispr', None, y_pred)
fr.store('pr1 pr2 pr3', None, y_pred_prob)
```

In Stata, we change the current working frame to **test**. We attach the value label **species** to **irispr**, and use the **tabulate** command to display a classification table. We also list the flowers that have been misclassified.

In [20]:

```

%%stata
frame change test
label values irispr species
label variable irispr predicted
tabulate iris irispr, row
list iris irispr pr1 pr2 pr3 if iris!=irispr

```

```

. frame change test

. label values irispr species

. label variable irispr predicted

. tabulate iris irispr, row

```

Key
frequency
row percentage

Iris species	predicted			Total
	setosa	versicolor	virginica	
setosa	11 100.00	0 0.00	0 0.00	11 100.00
versicolor	0 0.00	9 75.00	3 25.00	12 100.00
virginica	0 0.00	1 14.29	6 85.71	7 100.00
Total	11 36.67	10 33.33	9 30.00	30 100.00

```

. list iris irispr pr1 pr2 pr3 if iris!=irispr

```

	iris	irispr	pr1	pr2	pr3
16.	versicolor	virginica	.0036523	.3280099	.6683378
18.	versicolor	virginica	.0004762	.0687575	.9307663
19.	versicolor	virginica	.0031241	.4210142	.5758616
26.	virginica	versicolor	.0046717	.5435431	.4517851

Example 4: Integrate with Mata

This example will illustrate how to use the `%%mata` magic command to combine Python's capabilities with features of **Mata**, Stata's matrix programming language.

For illustrative purposes, we generate two random NumPy arrays in Python, **X** and **y**, representing the observations of the independent variables and the response variable. Our goal is to fit a linear regression and get the coefficients b , their standard errors, t statistics and p -values.

In [21]:

```
import numpy as np
np.random.seed(16)

# 100,000 observations and 7 variables
X = np.random.random((100000, 7))
y = np.random.random((100000, 1))
```

Then we push **X** and **y** to Mata using the magic command `%%mata`, and fit the model.

In [22]:

```

%%mata -m X, y
//add the constant term to X
cons = J(100000, 1, 1)
X = (X, cons)

// calculate b
b = invsym(X'X)*X'y

// calculate residuals and square of s
e = y - X*b
n = rows(X)
k = cols(X)
s2 = (e'e)/(n-k)

// calculate V
V = s2*invsym(X'X)

```

```

. mata
----- mata (type end to exit) -----
: //add the constant term to X
: cons = J(100000, 1, 1)

: X = (X, cons)

:
: // calculate b
: b = invsym(X'X)*X'y

:
: // calculate residuals and square of s
: e = y - X*b

: n = rows(X)

: k = cols(X)

: s2 = (e'e)/(n-k)

:
: // calculate V
: V = s2*invsym(X'X)

: end
-----

```

Next we calculate the standard errors, t statistics and p -values. Note that the standard errors are just the square roots of the diagonals of \mathbf{V} .

In [23]:

```

%%mata
se = sqrt(diagonal(V))
tstat = b:/se
pval = 2*ttail(n-k, abs(b:/se))

```

```

. mata
----- mata (type end to exit) -----
: se = sqrt(diagonal(V))

: tstat = b:/se

: pval = 2*ttail(n-k, abs(b:/se))

: end
-----

```

Next, we store the above results into a Mata matrix **ols_est** and push it back to Python.

In [24]:

```

%%mata -outm ols_est
ols_est = (b, se, tstat, pval)
ols_est

```

```

. mata
----- mata (type end to exit) -----
: ols_est = (b, se, tstat, pval)

: ols_est

```

	1	2	3	4
1	-.0028814142	.0031559598	-.9130072687	.36124092
2	.0021227247	.0031590697	.6719461302	.501619544
3	-.0001837604	.0031580164	-.058188558	.953598551
4	.0002157194	.0031606842	.068250868	.945586071
5	.001669965	.003153824	.5295048033	.5964564888
6	-.0009795437	.0031563296	-.3103426322	.7563010614
7	.0029998483	.003158516	.9497650963	.3422339202
8	.4988382533	.0042774646	116.6200762	0

```

: end
-----

```

In Python, we now have a NumPy array named **ols_est** with the same contents as the above Mata matrix.

In [25]:

ols_est

Out[25]:

```
array([[ -2.88141423e-03,  3.15595979e-03, -9.13007269e-01,
         3.61240920e-01],
       [ 2.12272468e-03,  3.15906973e-03,  6.71946130e-01,
         5.01619544e-01],
       [-1.83760422e-04,  3.15801643e-03, -5.81885580e-02,
         9.53598551e-01],
       [ 2.15719443e-04,  3.16068425e-03,  6.82508680e-02,
         9.45586071e-01],
       [ 1.66996498e-03,  3.15382404e-03,  5.29504803e-01,
         5.96456489e-01],
       [-9.79543651e-04,  3.15632965e-03, -3.10342632e-01,
         7.56301061e-01],
       [ 2.99984826e-03,  3.15851600e-03,  9.49765096e-01,
         3.42233920e-01],
       [ 4.98838253e-01,  4.27746465e-03,  1.16620076e+02,
         0.00000000e+00]])
```

We can compare **ols_est** with the results reported by Stata's **regress** command. Below, we will use the **%%stata** magic command to execute it. Before we do that, we combine **X** and **y** into one NumPy array called **data**.

In [26]:

```
data = np.concatenate((X, y), axis=1)
```

Then we push **data** to Stata. When reading a NumPy array into Stata, the variables are named **v1**, **v2**, ... by default, so we rename them. Then we fit the regression model.

In [27]:

```
%%stata -d data -force
describe

// rename the last variable to y
rename v8 y

// rename the other variables, prefixing them with an x
rename v* x*

regress y x1-x7
```

```
. describe
```

Contains data

```
obs:      100,000
vars:      8
```

variable name	storage type	display format	value label	variable label
v1	double	%10.0g		
v2	double	%10.0g		
v3	double	%10.0g		
v4	double	%10.0g		
v5	double	%10.0g		
v6	double	%10.0g		
v7	double	%10.0g		
v8	double	%10.0g		

Sorted by:

Note: Dataset has changed since last saved.

```
.
. // rename the last variable to y
. rename v8 y

.
. // rename the other variables, prefixing them with an x
. rename v* x*

.
. regress y x1-x7
```

Source	SS	df	MS	Number of obs	=	100,000
				F(7, 99992)	=	0.37
Model	.213224988	7	.030460713	Prob > F	=	0.9218
Residual	8297.31719	99,992	.08297981	R-squared	=	0.0000
				Adj R-squared	=	-0.0000
Total	8297.53042	99,999	.082976134	Root MSE	=	.28806

y	Coef.	Std. Err.	t	P> t	[95% Conf. Interval]	
x1	-.0028814	.003156	-0.91	0.361	-.0090671	.0033042
x2	.0021227	.0031591	0.67	0.502	-.004069	.0083145
x3	-.0001838	.003158	-0.06	0.954	-.0063734	.0060059
x4	.0002157	.0031607	0.07	0.946	-.0059792	.0064106
x5	.00167	.0031538	0.53	0.596	-.0045115	.0078514
x6	-.0009795	.0031563	-0.31	0.756	-.0071659	.0052068
x7	.0029998	.0031585	0.95	0.342	-.0031908	.0091905
_cons	.4988383	.0042775	116.62	0.000	.4904545	.507222

Example 5: Call Stata using API functions

In addition to the magic commands, you can also call Stata using the API functions defined in the **stata** module. For illustration purpose, we use the Boston housing price dataset from the **scikit-learn** package.

In [28]:

```
from sklearn import datasets  
  
bos = datasets.load_boston()  
print(bos.DESCR)
```

```
.. _boston_dataset:
```

Boston house prices dataset

****Data Set Characteristics:****

:Number of Instances: 506

:Number of Attributes: 13 numeric/categorical predictive. Median Value (attribute 14) is usually the target.

:Attribute Information (in order):

- CRIM	per capita crime rate by town
- ZN	proportion of residential land zoned for lots over 25,000 sq. ft.
- INDUS	proportion of non-retail business acres per town
- CHAS	Charles River dummy variable (= 1 if tract bounds river; 0 otherwise)
- NOX	nitric oxides concentration (parts per 10 million)
- RM	average number of rooms per dwelling
- AGE	proportion of owner-occupied units built prior to 1940
- DIS	weighted distances to five Boston employment centres
- RAD	index of accessibility to radial highways
- TAX	full-value property-tax rate per \$10,000
- PTRATIO	pupil-teacher ratio by town
- B	$1000(B_k - 0.63)^2$ where B_k is the proportion of blacks by town
- LSTAT	% lower status of the population
- MEDV	Median value of owner-occupied homes in \$1000's

:Missing Attribute Values: None

:Creator: Harrison, D. and Rubinfeld, D.L.

This is a copy of UCI ML housing dataset.

<https://archive.ics.uci.edu/ml/machine-learning-databases/housing/>

This dataset was taken from the StatLib library which is maintained at Carnegie Mellon University.

The Boston house-price data of Harrison, D. and Rubinfeld, D.L. 'Hedonic prices and the demand for clean air', J. Environ. Economics & Management, vol.5, 81-102, 1978. Used in Belsley, Kuh & Welsch, 'Regression diagnostics ...', Wiley, 1980. N.B. Various transformations are used in the table on pages 244-261 of the latter.

The Boston house-price data has been used in many machine learning papers that address regression problems.

.. topic:: References

- Belsley, Kuh & Welsch, 'Regression diagnostics: Identifying Influential Data and Sources of Collinearity', Wiley, 1980. 244-261.
- Quinlan, R. (1993). Combining Instance-Based and Model-Based Learning. In Proceedings on the Tenth International Conference of Machine Learning, 236-243, University of Massachusetts, Amherst. Morgan Kaufmann.

We store the dataset into a Pandas dataframe named **boston**.

In [29]:

```
import pandas as pd
boston = pd.DataFrame(bos.data)
boston.columns = bos.feature_names
boston['MEDV'] = bos.target
boston.head()
```

Out[29]:

	CRIM	ZN	INDUS	CHAS	NOX	RM	AGE	DIS	RAD	TAX	PTRATIO	B
0	0.00632	18.0	2.31	0.0	0.538	6.575	65.2	4.0900	1.0	296.0	15.3	396.90
1	0.02731	0.0	7.07	0.0	0.469	6.421	78.9	4.9671	2.0	242.0	17.8	396.90
2	0.02729	0.0	7.07	0.0	0.469	7.185	61.1	4.9671	2.0	242.0	17.8	392.83
3	0.03237	0.0	2.18	0.0	0.458	6.998	45.8	6.0622	3.0	222.0	18.7	394.63
4	0.06905	0.0	2.18	0.0	0.458	7.147	54.2	6.0622	3.0	222.0	18.7	396.90

Then we load the Pandas dataframe into Stata using the **dataframe_to_data()** function of the **stata** module.

In [30]:

```
from pystata import stata
stata.dataframe_to_data(boston, force=True)
```

Next, we specify the variable labels and attach a value label to **CHAS**, which indicates whether the tract bounds the Charles River, using the **run()** method.

In [31]:

```
stata.run('''
label variable MEDV "Median value of owner-occupied homes in $1000s"
label variable CRIM "Town crime rate, per capita"
label variable RM "Average number of rooms per dwelling"
label variable CHAS "Tract bounds the Charles River (= 1 if tract bounds river; 0 otherwise)"

label define bound 1 "Yes" 0 "No", replace
label values CHAS bound
''')
```

```
.
. label variable MEDV "Median value of owner-occupied homes in $1000s"

. label variable CRIM "Town crime rate, per capita"

. label variable RM "Average number of rooms per dwelling"

. label variable CHAS "Tract bounds the Charles River (= 1 if tract bounds rive
> r; 0 otherwise)"

.
. label define bound 1 "Yes" 0 "No", replace

. label values CHAS bound
```

Afterwards, we fit a linear regression model of the median home value (**MEDV**) on the town crime rate (**CRIM**), the average number of rooms per dwelling (**RM**), and whether the house is close to the Charles River (**CHAS**). Then we predict the median home values, storing these values in the variable **midval**.

In [32]:

```
stata.run('''
regress MEDV CRIM RM i.CHAS
predict midval
''')
```

```
. regress MEDV CRIM RM i.CHAS
```

Source	SS	df	MS	Number of obs	=	506
				F(3, 502)	=	206.73
Model	23607.3566	3	7869.11888	Prob > F	=	0.0000
Residual	19108.9388	502	38.0656151	R-squared	=	0.5527
				Adj R-squared	=	0.5500
Total	42716.2954	505	84.5867236	Root MSE	=	6.1697

MEDV	Coef.	Std. Err.	t	P> t	[95% Conf. Interval]
CRIM	-.2607244	.0327369	-7.96	0.000	-.3250427 -.1964061
RM	8.27818	.4018204	20.60	0.000	7.488723 9.067637
CHAS					
Yes	3.763037	1.086199	3.46	0.001	1.628981 5.897093
_cons	-28.81068	2.563308	-11.24	0.000	-33.84682 -23.77455

```
. predict midval
(option xb assumed; fitted values)
```

Then we use the **get_ereurn()** function to push all the **e()** results to Python as a dictionary named **steret**. And we use the **data_to_array()** function to store the predicted values (**midval**) in a NumPy array named **stpred**.

In [33]:

```
steret = stata.get_ereurn()
stpred = stata.data_to_array('midval')
```

Conclusion

- The **pystata** package provides tight integration bewtween Stata and Python using IPython magic commands and a suite of API functions.
- Users can execute Stata code in all Python's environment to leverage Stata's capabilities.
- The **pystata** package can be used together with the [Stata Function Interface \(sfi\)](https://www.stata.com/python/api16/index.html) (<https://www.stata.com/python/api16/index.html>) module to access Stata's core features.